

```
In [52]: import os  
os.getcwd()
```

```
Out[52]: '/home/sapatevaibhav/Documents/ML'
```

```
In [1]: import pandas as pd
```

```
In [5]: df = pd.read_csv('SMSSpamCollection', sep='\t', names=['tag', 'data']);
```

```
In [6]: df
```

```
Out[6]:
```

	tag	data
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
...
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will ü b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

5572 rows × 2 columns

```
In [9]: import nltk  
nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to  
[nltk_data] /home/sapatevaibhav/nltk_data...  
[nltk_data] Unzipping corpora/stopwords.zip.
```

```
Out[9]: True
```

```
In [15]: from nltk.corpus import stopwords  
swords = stopwords.words('english')  
from nltk.stem import PorterStemmer  
ps = PorterStemmer()
```

```
In [21]: from sklearn.feature_extraction.text import TfidfVectorizer  
from nltk.tokenize import word_tokenize
```

```
In [22]: def clean_text(sent):
          tokens = word_tokenize(sent)
          clean = [word for word in tokens
                   if word.isdigit() or word.isalpha()]
          clean = [ps.stem(word) for word in clean
                   if word not in swords]
          return clean
```

```
In [17]: tfidf = TfidfVectorizer(analyzer = clean_text)
```

```
In [25]: nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to
[nltk_data]   /home/sapatevaibhav/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
```

```
Out[25]: True
```

```
In [32]: x = df['data']
          y = df['tag']
          x_new = tfidf.fit_transform(x)
```

```
In [28]: x.shape
```

```
Out[28]: (5572,)
```

```
In [29]: x_new.shape
```

```
Out[29]: (5572, 6513)
```

```
In [30]: x_new
```

```
Out[30]: <5572x6513 sparse matrix of type '<class 'numpy.float64''>'
          with 52578 stored elements in Compressed Sparse Row format>
```

```
In [31]: y.value_counts()
```

```
Out[31]: ham      4825
          spam      747
          Name: tag, dtype: int64
```

```
In [34]: from sklearn.model_selection import train_test_split
          x_train, x_test, y_train, y_test = train_test_split(x_new, y, random_state = 0,
```


```
In [35]: from sklearn.naive_bayes import GaussianNB
```

```
In [37]: nb = GaussianNB()
```

```
In [38]: nb.fit(x_train.toarray(), y_train)
```

```
Out[38]: ▼ GaussianNB
          GaussianNB()
```

```
In [39]: y_pred = nb.predict(x_test.toarray())
          from sklearn.metrics import ConfusionMatrixDisplay
          ConfusionMatrixDisplay.from_predictions(y_test, y_pred);
```



```
In [41]: from sklearn.metrics import accuracy_score, classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
ham	0.98	0.87	0.92	1208
spam	0.51	0.89	0.65	185
accuracy			0.87	1393
macro avg	0.75	0.88	0.79	1393
weighted avg	0.92	0.87	0.89	1393


```
In [42]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(random_state = 0)
```

```
In [43]: rf.fit(x_train, y_train)
```

```
Out[43]: ▼      RandomForestClassifier
RandomForestClassifier(random_state=0)
```

```
In [44]: y_pred = rf.predict(x_test)
ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
```

```
Out[44]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fe716527d60>
```



```
In [45]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
ham	0.98	1.00	0.99	1208
spam	0.99	0.86	0.92	185
accuracy			0.98	1393
macro avg	0.98	0.93	0.96	1393
weighted avg	0.98	0.98	0.98	1393

```
In [46]: from sklearn.linear_model import LogisticRegression
log = LogisticRegression()
log.fit(x_train, y_train)
y_pred = log.predict(x_test)
accuracy_score(y_test, y_pred)
```

```
Out[46]: 0.9641062455132807
```

```
In [47]: from sklearn.model_selection import GridSearchCV
params = {
    'criterion': ['gini', 'entropy'],
    'max_features': ['sqrt', 'log2'],
    'random_state': [0, 1, 2, 3, 4],
    'class_weight': ['balanced', 'balanced_subsample']
}
```

```
In [48]: grid = GridSearchCV(rf, param_grid = params, cv = 5, scoring = 'accuracy')
```

```
In [49]: grid.fit(x_train, y_train)
```

Out[49]:

► **GridSearchCV**

► **estimator: RandomForestClassifier**

► RandomForestClassifier

In [50]: `rf = grid.best_estimator_`

In [51]: `y_pred = rf.predict(x_test)`
`accuracy_score(y_test, y_pred)`

Out[51]: 0.9777458722182341